

What is a Neuron (in Deep Learning)?

A **neuron** in deep learning is a **computational unit** that:

1. Takes inputs
2. Applies weights
3. Adds bias
4. Passes the result through an **activation function**
5. Produces an output

Just like a biological neuron processes signals, this **artificial neuron** processes numbers.

Formula of a Neuron:

Output= Activation($w_1x_1+w_2x_2+\dots+w_nx_n+b$)

Where:

- x_1, x_2, \dots, x_n : Inputs
 - w_1, w_2, \dots, w_n : Weights
 - b : Bias
 - Activation: Function like sigmoid, ReLU
-

Example: Let's say you have a single neuron

Inputs:

- $x_1=2, x_2=3$

Weights:

- $w_1=0.4, w_2=0.6$

Bias:

- $b=0.5$

Activation Function:

- Sigmoid:

$$\sigma(z) = 1/(1+e^{-z})$$

Step-by-Step Calculation:

1. Weighted sum:

$$Z = (0.4 \times 2) + (0.6 \times 3) + 0.5 = 0.8 + 1.8 + 0.5 = 3.1$$

2. Apply activation (sigmoid):

Output = $1 / (1 + e^{-3.1}) \approx 0.957$

✅ The neuron **outputs 0.957**.

Purpose of Neuron:

Each neuron **extracts patterns** from the data. In a network:

- First layer might detect simple features (edges in image)
 - Next layers detect complex features (shapes, objects)
-

Real-World Analogy:

Imagine a **restaurant chef**:

- Inputs: Ingredients (like flour, veggies)
 - Weights: How much of each ingredient
 - Bias: A secret spice for adjustment
 - Activation: Cooking method (bake/fry)
 - Output: Final dish (prediction)
-

In a Neural Network:

A **neuron is one brick**.

Many neurons = one layer

Many layers = a deep network

What is a Perceptron?

A **perceptron** is the **simplest type of neural network**.

It consists of a **single neuron** that:

- Takes **multiple inputs**
- Applies **weights** and **bias**
- Passes the result through an **activation function**
- Produces **binary output** (like 0 or 1)

Mathematical Formula:

$$\text{Output} = \begin{cases} 1 & \text{if } w_1x_1 + w_2x_2 + \dots + w_nx_n + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

This is also known as a **threshold function** (like a yes/no decision).

Purpose:

The perceptron is used for **binary classification** — like:

- Email: spam or not spam
 - Image: cat or not cat
 - Disease: yes or no
-

Simple Manual Example:

Let's say we want to build a perceptron that decides whether a student **passes** or **fails** based on:

- x_1 = hours studied
- x_2 = number of classes attended

Let's assume:

- $w_1 = 0.6$
- $w_2 = 0.5$
- $b = -0.7$

Now we give an input:

- $x_1 = 1$ (studied 1 hour)
- $x_2 = 1$ (attended 1 class)

Step-by-step:

$$Z = w_1x_1 + w_2x_2 + b$$

$$= (0.6)(1) + (0.5)(1) - 0.7$$

$$= 0.4$$

Since $z > 0$,

Output = 1 → The student **passes**.

Perceptron Learning Rule (Weight Update):

If the output is wrong:

$$W_i = w_i + \Delta w_i$$

$$\Delta w_i = \eta (y - y^{\wedge}) x_i$$

Where:

- η = learning rate
 - y = true label
 - y^{\wedge} = predicted output
 - x_i = input value
-

Key Points:

Term	Meaning
------	---------

Inputs (x)	Features of the data
------------	----------------------

Weights (w)	Importance of each feature
-------------	----------------------------

Bias (b)	Shifts the decision boundary
----------	------------------------------

Activation	Step function (output 0 or 1)
------------	-------------------------------

Output	Prediction based on threshold
--------	-------------------------------

Limitation:

- A single-layer perceptron can **only solve linearly separable problems**.
- Can't solve **XOR problem**, for example.

That's why **multi-layer perceptron (MLPs)** or **deep networks** are used in modern deep learning.

Real-World Analogy:

Imagine a **gatekeeper** checking two things:

- Did the person bring a ticket?
- Is the person on time?

If both are true (weighted sum exceeds a threshold), entry is allowed. Else, not.

Weights and Bias — in simple language with examples.

1. What is a Weight (w)?

A **weight** is a number that **determines how important** a particular input is to the neuron's output.

Think of it like:

"How much attention should the neuron pay to this input?"

- A **high positive weight** means strong influence.
 - A **negative weight** means the input has a **negative effect**.
 - A weight of **zero** means the input is **ignored**.
-

2. What is a Bias (b)?

Bias is a **constant value added** to the weighted sum of inputs **before** applying the activation function.

Think of it like:

"Even if all inputs are 0, should the neuron still be activated?"

Bias **shifts the output** of the activation function — helping the model **fit the data better**.

Neuron Output Formula:

Output = Activation ($w_1x_1 + w_2x_2 + \dots + w_nx_n + b$)

Where:

- x_i : Inputs
 - w_i : Weights
 - b : Bias
 - Activation: Sigmoid, ReLU, etc.
-

Simple Example:

Suppose you're predicting **whether a student passes or fails**, based on:

- x_1 : hours studied
- x_2 : hours of sleep

Let's use:

- $w_1 = 0.6$ $w_2 = 0.4$, $b = -0.7$

Inputs:

- $x_1 = 2$, $x_2 = 5$

Now compute:

$$Z = (0.6)(2) + (0.4)(5) - 0.7 = 1.2 + 2.0 - 0.7 = 2.5$$

Output = Activation(2.5)

✅ Because the weighted sum is high → neuron activates → student passes.

Analogy:

Imagine you're judging whether to go for a **trip**.

- x_1 : weather (sunny = 1, rainy = 0)
 - x_2 : friends joining (yes = 1, no = 0)
 - **Weights** tell how much you care about each.
 - weather: weight = 0.8
 - friends: weight = 0.2
 - **Bias** is your personal preference.
 - Even if weather and friends are bad, you still want to go? That's bias!
-

Why are Weights and Bias Important?

Component Purpose

Weight Learns the **importance of each input** during training

Bias Helps the model **fit better by shifting activation** (adds flexibility)

Summary:

- Weights are like **sliders** that adjust how much each input matters.
- Bias is like a **threshold adjuster** that moves the decision boundary.
- Both are **learned during training** using optimization techniques like **gradient descent**.

What is an Activation Function?

An **activation function** is a mathematical function used in a **neuron** of a neural network to:

1. **Decide whether the neuron should be activated (fire)** or not
2. **Introduce non-linearity** into the model (which is crucial for learning complex patterns)

Real-Life Analogy:

Think of it like a **decision switch**.

If inputs (weighted sum) are good enough, activate the neuron.
Else, keep it silent.

Formula Context:

In a neural network:

$$Z = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

$$\text{Output} = \text{Activation}(z)$$

So, the activation function is applied **after the weighted sum and bias**.

Why Do We Need Activation Functions?

Without an activation function:

- The neural network is just doing **linear calculations**.
- Can't solve **complex tasks** like image recognition, language translation, etc.

With activation:

- ✓ The model becomes **non-linear** and can **learn complex patterns**.

Common Activation Functions:

Function	Formula / Shape	Use Case Example
Step Function	Output: 0 or 1 (binary)	Simple binary classification (Perceptron)
Sigmoid	$\sigma(z) = 1 / (1 + e^{-z})$	Probabilities in binary classification
Tanh	$\text{Tanh}(z) = (e^z - e^{-z}) / (e^z + e^{-z})$	Better zero-centered outputs
ReLU	$\text{ReLU}(z) = \max(0, z)$	Deep learning — fast, efficient, widely used
Leaky ReLU	$\max(0.01z, z)$	Solves dead neuron problem in ReLU

Function	Formula / Shape	Use Case Example
Softmax	Converts logits to probabilities	Multi-class classification (last layer)

Example with ReLU:

Suppose:

- Weighted sum = $z = -3$

$$\text{ReLU}(-3) = \max(0, -3) = 0$$

Now, if:

- $z=5$

$$\text{ReLU}(5) = \max(0, 5) = 5$$

So it "lets through" only **positive values**, and stops negative ones.

Visual Intuition:

- **Sigmoid** looks like an "S" curve
 - **ReLU** is like a ramp (flat at 0, linear after 0)
 - **Softmax** converts scores into percentages (adds up to 100%)
-

Summary:

Feature	Description
Adds non-linearity	Enables network to learn complex patterns
Determines Firing	Controls whether neuron outputs signal or not
Helps in Classification	Maps values to probabilities in some cases

Forward Propagation and Backward Propagation

Step-by-Step Overview

We'll build a tiny neural network:

- **1 input layer (2 neurons)**
 - **1 hidden layer (2 neurons)**
 - **1 output layer (1 neuron)**
 - Use **Sigmoid** as the activation function
 - We'll do both **Forward Propagation** and then **Backward Propagation** to **calculate gradients**.
-

Let's Define the Network

Inputs:

- $x_1 = 0.05$
- $x_2 = 0.10$

Target Output:

- $y = 0.01$

Initial Weights:

Input to Hidden Weights:

	Hidden1 (h1)	Hidden2 (h2)
x1	0.15	0.25
x2	0.20	0.30

Bias to Hidden:

- $b_{h1} = b_{h2} = 0.35$

Hidden to Output Weight:

	Output (o1)
h1	0.40
h2	0.45

Bias to Output:

- $bo = 0.60$
-

Forward Propagation

Step 1: Hidden Layer Inputs

$$neth1 = x1 \cdot w1 + x2 \cdot w2 + b_{h1}$$

$$= (0.05)(0.15) + (0.10)(0.20) + 0.35$$

$$= 0.0075 + 0.020 + 0.35$$

$$= 0.3775$$

$$neth2 = (0.05)(0.25) + (0.10)(0.30) + 0.35$$

$$= 0.0125 + 0.030 + 0.35$$

$$= 0.3925$$

Step 2: Activation (Sigmoid)

$$outh1 = \sigma(0.3775)$$

$$= 1 / (1 + e^{-0.3775}) \approx 0.59327$$

$$outh2 = \sigma(0.3925)$$

$$= 1 / (1 + e^{-0.3925}) \approx 0.59688$$

Step 3: Output Layer Input

$$neto1 = outh1 \cdot w5 + outh2 \cdot w6 + bo$$

$$= (0.59327)(0.40) + (0.59688)(0.45) + 0.60$$

$$= 0.2373 + 0.2686 + 0.60$$

$$= 1.1059$$

Step 4: Output Layer Activation

$$outo1 = \sigma(1.1059)$$

$$= 1 / (1 + e^{-1.1059}) \approx 0.75136$$

Step 5: Calculate Error (Loss)

Using Mean Squared Error:

$$E = 12(\text{target} - \text{output})^2$$

$$= 12 (0.01 - 0.75136)^2 \approx 0.2748$$

What is Backward Propagation?

Backward Propagation (or **Backprop**) is the process of calculating the **gradient of the loss function** with respect to each weight and bias, and **updating them** using **gradient descent** to minimize the error.

We use the **chain rule from calculus** to break down how the error changes with respect to each parameter.

Network Structure (Reminder)

We'll use the **same small neural network**:

- Inputs: $x_1 = 0.05$, $x_2 = 0.10$
- Target: $y = 0.01$
- One hidden layer (2 neurons), one output neuron
- Activation: **Sigmoid**
- Loss Function: **Mean Squared Error**

Initial weights:

Input → Hidden:

$$w_1 = 0.15, w_2 = 0.20$$

$$w_3 = 0.25, w_4 = 0.30$$

$$\text{Bias for hidden} = 0.35$$

Hidden → Output:

$$w_5 = 0.40, w_6 = 0.45$$

$$\text{Bias for output} = 0.60$$

Goal of Backpropagation

We want to **update all weights (w_1 to w_6)** to **reduce the error**.

This requires:

1. Compute the gradient of error w.r.t each weight
2. Update weights:

$$W = w - \eta \cdot \partial E / \partial w$$

Where η (eta) is the learning rate, say 0.5.

Step-by-Step: Backward Propagation

Let's break this into 2 parts:

PART A: Update weights from Hidden → Output

We're updating **w5** and **w6**, which connect h1 and h2 to output o1.

Let's recall:

- Predicted output out_o1 = 0.75136
 - Target y = 0.01
 - Output from hidden neurons:
out_h1 = 0.59327, out_h2 = 0.59688
-

Step A1: Compute $\partial E / \partial w5$ using Chain Rule

We use:

$$\partial E / \partial w5 = \partial E / \partial \text{outo1} * \partial \text{outo1} / \partial \text{neto1} * \partial \text{neto1} / \partial w5$$

Breakdown:

1. Error derivative:

$$\partial E / \partial \text{outo1} = \text{outo1} - y = 0.75136 - 0.01 = 0.74136$$

2. Derivative of sigmoid:

$$\partial \text{outo1} / \partial \text{neto1} = \text{outo1} * (1 - \text{outo1}) = 0.75136 * 0.24864 \approx 0.1868$$

3. Net to weight derivative:

$$\partial \text{neto1} / \partial w5 = \text{outh1} = 0.59327$$

Now:

$$\partial E / \partial w5 = 0.74136 * 0.1868 * 0.59327 \approx 0.08216$$

Similarly:

$$\partial E / \partial w6 = 0.74136 * 0.1868 * 0.59688 \approx 0.08267$$

◆ Step A2: Update Weights

$$w5 = w5 - \eta * \partial E / \partial w5$$

$$= 0.40 - 0.5 * 0.08216 \approx 0.35892$$

$$w6 = 0.45 - 0.5 * 0.08267 \approx 0.40867$$

✓ **Output layer weights updated!**

■ PART B: Update weights from Input → Hidden

Now we go **deeper** to adjust $w1, w2, w3, w4$.

Let's do for **$w1$** (from $x1 \rightarrow h1$).

◆ Step B1: Use Chain Rule for $w1$

$$\partial E / \partial w1 = \partial E / \partial \text{outo1} * \partial \text{outo1} / \partial \text{neto1} * \partial \text{neto1} / \partial \text{outh1} * \partial \text{outh1} / \partial \text{neth1} * \partial \text{neth1} / \partial w1$$

- $\partial E / \partial \text{outo1} = 0.74136$
- $\partial \text{outo1} / \partial \text{neto1} = 0.1868$
- $\partial \text{neto1} / \partial \text{outh1} = w5 = 0.40$

Now compute:

- $\partial \text{outh1} / \partial \text{neth1} = \text{outh1}(1 - \text{outh1})$
 $= 0.59327(1 - 0.59327) \approx 0.2413$
- $\partial \text{neth1} / \partial w1 = x1 = 0.05$

Now multiply:

$$\partial E / \partial w1 = 0.74136 * 0.1868 * 0.40 * 0.2413 * 0.05 \approx 0.00067$$

◆ Step B2: Update $w1$

$$w1 = 0.15 - 0.5 * 0.00067$$

$$= 0.14967$$

Repeat this for $w2, w3$, and $w4$ in the same way.

🧠 What Did We Learn?

- Backpropagation is all about using **calculus chain rule** to move error from output → hidden → input.
- Every weight is updated by seeing **how much it contributed to the final error**.
- The smaller the contribution, the smaller the update.

- This is repeated for **many epochs**.
-

📌 Summary Table

Step	Explanation
$\partial E / \partial w_5$	Error at output layer and how it flows back
Chain Rule	To connect the effect of weight on total error
Update Rule	Gradient Descent: New weight = old - $\eta \times$ gradient
Repeat	Do this for all weights, layer by layer

🧠 What is an Optimizer in Neural Networks?

An **optimizer** is an algorithm that **adjusts the weights and biases** of a neural network to **reduce the loss (error)** during training.

It does this by:

- Using **gradients** calculated during **backpropagation**.
 - Deciding **how much and in which direction** each weight and bias should be updated.
-

📌 Why Optimizer is Important?

Without an optimizer:

- Your neural network won't learn anything.
 - You won't know **how to update weights** after calculating the error.
-

🎯 Goal of an Optimizer

Minimize the **loss function** (for example, MSE or Cross-Entropy):

$\min_{\theta} \text{Loss}(y_{\text{true}}, y_{\text{predicted}})$

Where θ represents all the weights and biases in the network.

How it Works (Conceptual Steps)

1. Start with **random weights**.
 2. Use **forward propagation** to compute predictions.
 3. Calculate **loss** (how wrong the prediction is).
 4. Use **backpropagation** to get gradients.
 5. The **optimizer** uses those gradients to **update weights** to reduce the error.
 6. Repeat steps 2–5 until convergence (or fixed epochs).
-

Common Optimizers

Optimizer	Key Idea	Good For
SGD	Update weights using the gradient only	Small/simple models
Momentum	Uses past updates to accelerate learning	Faster convergence
RMSprop	Scales learning rate for each parameter	RNNs, unstable gradients
Adam	Combines Momentum + RMSprop	Most deep learning models

Example: Using SGD (Stochastic Gradient Descent)

Suppose we are training a network with just **1 weight** $w = 0.5$, learning rate $\eta = 0.1$, and we got gradient of loss:

$$\partial E / \partial w = 0.4$$

Then optimizer (SGD) updates:



$$w_{\text{new}} = w - \eta \cdot \partial E / \partial w$$




$$= 0.5 - 0.1 \cdot 0.4$$

$$= 0.46$$

So weight becomes 0.46. The optimizer does this **for all weights** in the network during training.

Summary

Feature	Explanation
 Goal	Minimize the loss by adjusting weights
 How	Use gradients to take steps towards lower error

Feature	Explanation
 Step	Each training step updates weights using optimizer logic
 Examples	SGD, Adam, RMSprop, Momentum
 Best Practice	Use Adam for most use cases unless reason to switch

What is an Epoch in Deep Learning?

An **epoch** is **one complete pass** through the **entire training dataset** by the neural network.

In simple words:

If you have 1000 training samples, and your model sees **all 1000 once**, it has completed **1 epoch**.

Why Do We Use Multiple Epochs?

A single pass (1 epoch) is usually **not enough** for the model to learn the underlying patterns. So we train for **multiple epochs**, allowing the model to:

- Gradually **reduce the error**
 - Improve accuracy by adjusting weights through **backpropagation**
-

Related Concepts:

1. Batch size:

Number of samples the model sees **before updating weights**.

Example:

- Dataset: 1000 samples
- Batch size: 100
- \Rightarrow 10 batches per epoch

2. Iteration:

One **weight update** = 1 iteration

So:

Iterations per epoch = Total samples / Batch size

Example:

Assume you have:

```
X_train.shape = (1000, 10)
```

```
batch_size = 100
```

```
epochs = 10
```

Result:

10 batches per epoch

100 total batches = 10 epochs × 10 batches

As epochs increase, ideally:

- **Loss ↓**
 - **Accuracy ↑**
-

⚠ Too Many Epochs?

- If too few → **Underfitting**
- If too many → **Overfitting**

Use **Early Stopping** to avoid training too long.

📌 Summary

Term	Meaning
Epoch	1 full pass-through training data
Batch Size	Number of samples seen before update
Iteration	One update of weights
Goal	Train the model gradually over epochs